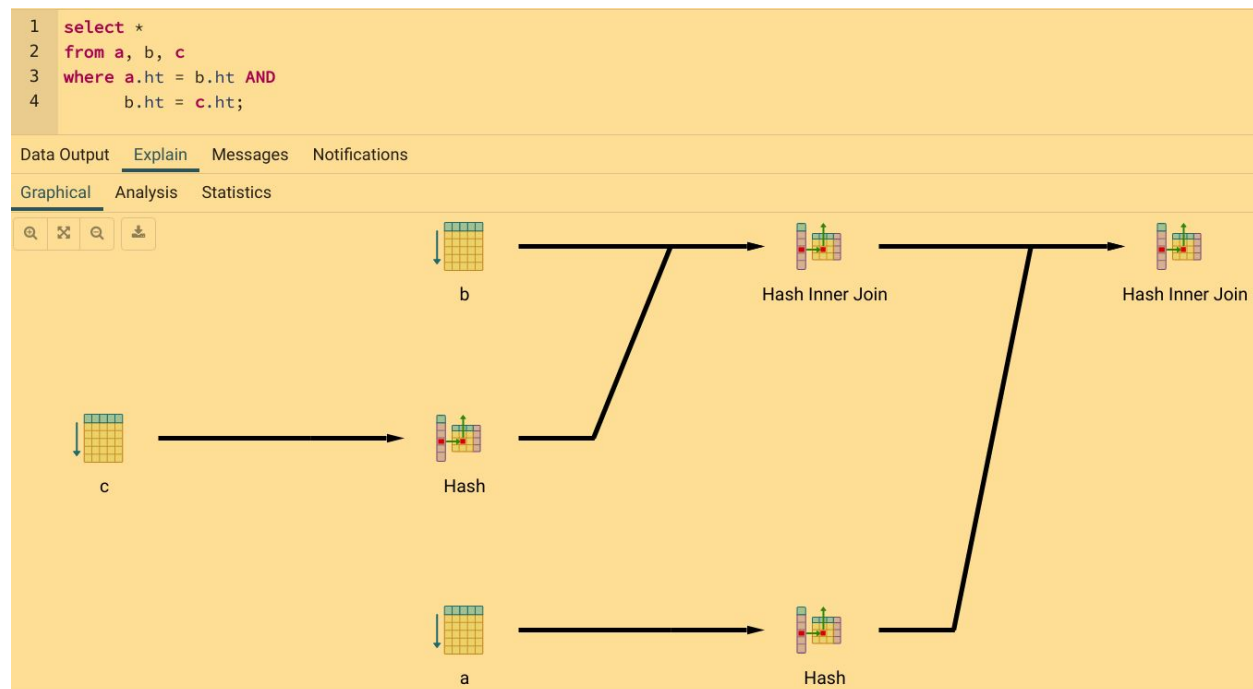


Part 1:

Client Location	Server Locaiton	Load Time	Bytes/second loaded
Local machine	Local machine	207s	6.18MB/s
Local machine	AWS	1551s	0.82MB/s
AWS	AWS	467.5s	2.74MB/s

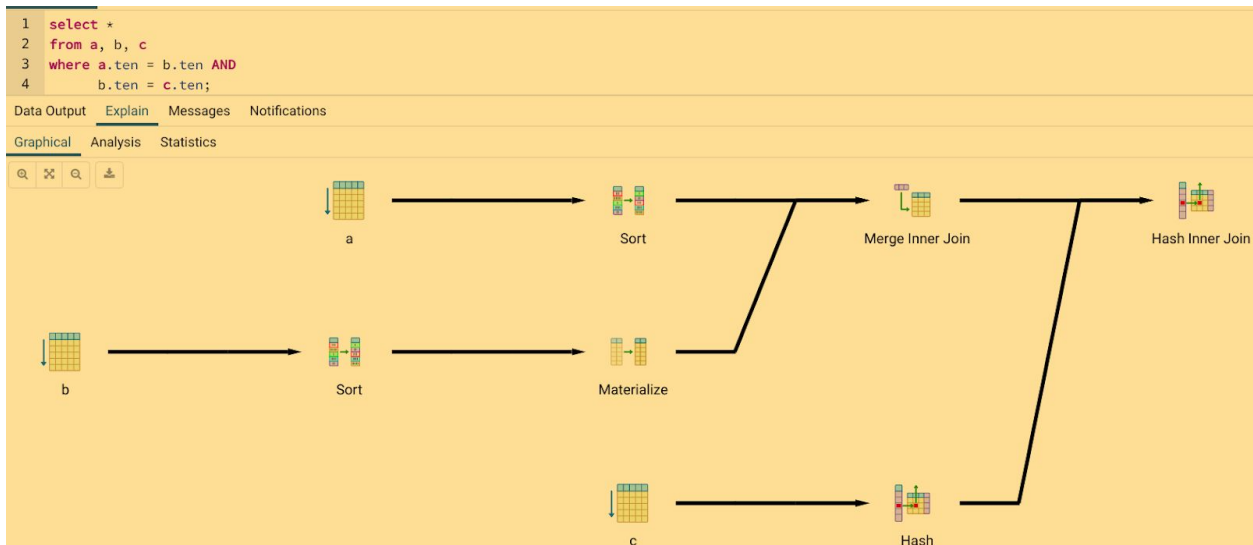
Part 2:

10.



The hash-join algorithm is chosen.

11.



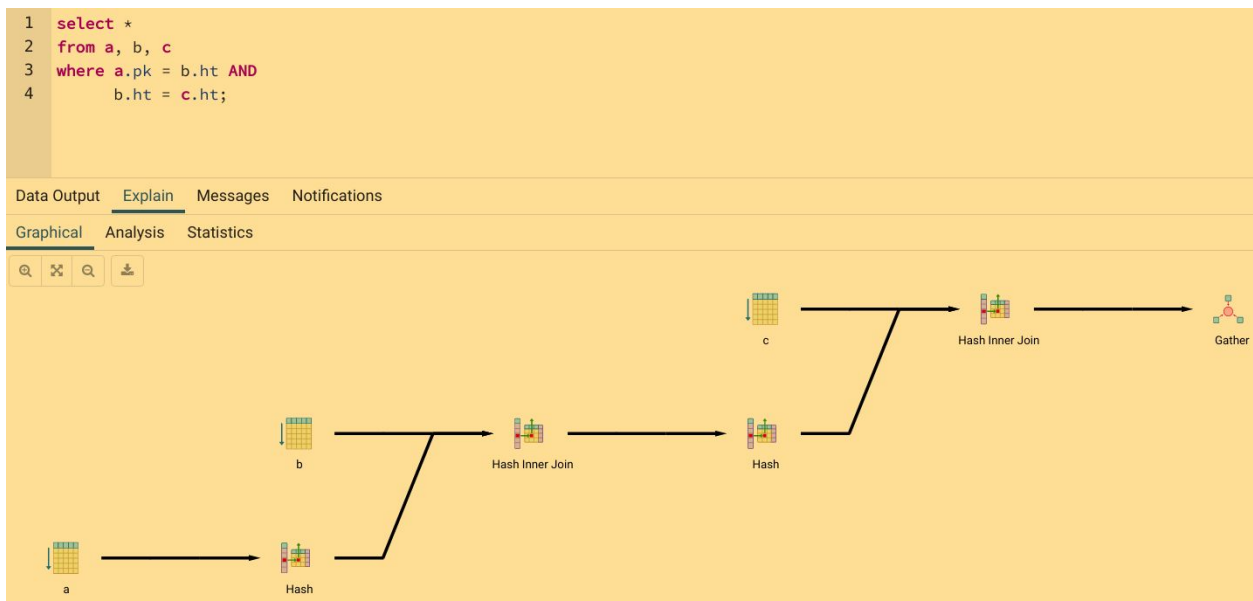
Surprisingly, when we join on the “ten” column which has a higher selectivity, sort and merge join is selected first to join a and b. Then the result is hash-joined with the third table c.

12.

When joining on “ten” or “hund” column, it’s using the query plan in 11;

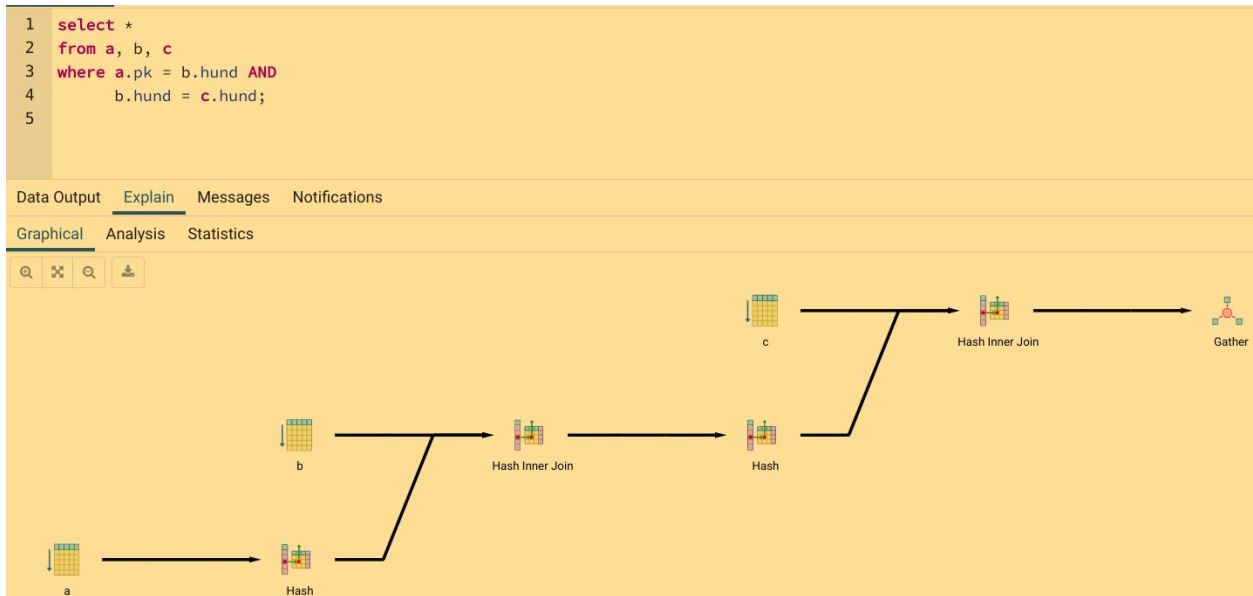
When joining on “ot”, “tt”, or “ht” column, it’s using the query plan in 10.

13.



This time, again, hash join is chosen and used twice to get the final result.

14.

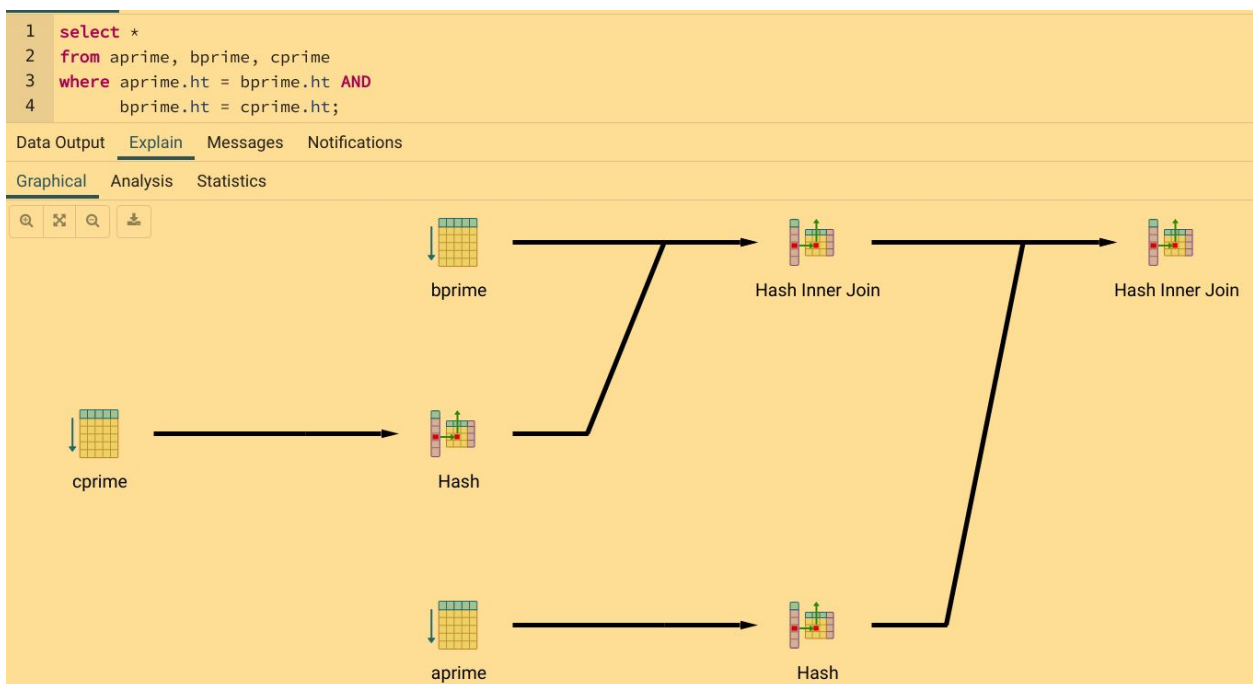


Similarly, hash join is used.

15.

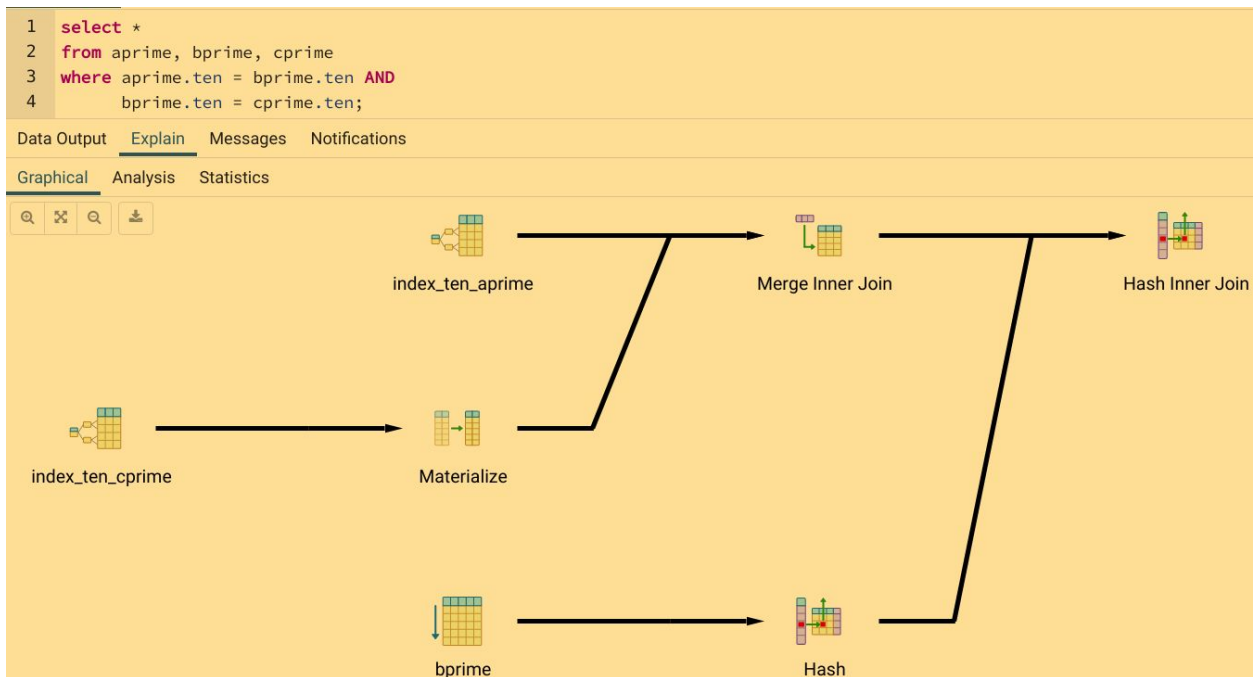
I tried to `A.pk = b.ten = c.ten` and it gives the same query plan! In other words, 13 and 14 produce the same query plan. Moreover, it's always going to be the same plan no matter which column we choose to join on b and c.

16.



Same, hash join is chosen.

17.

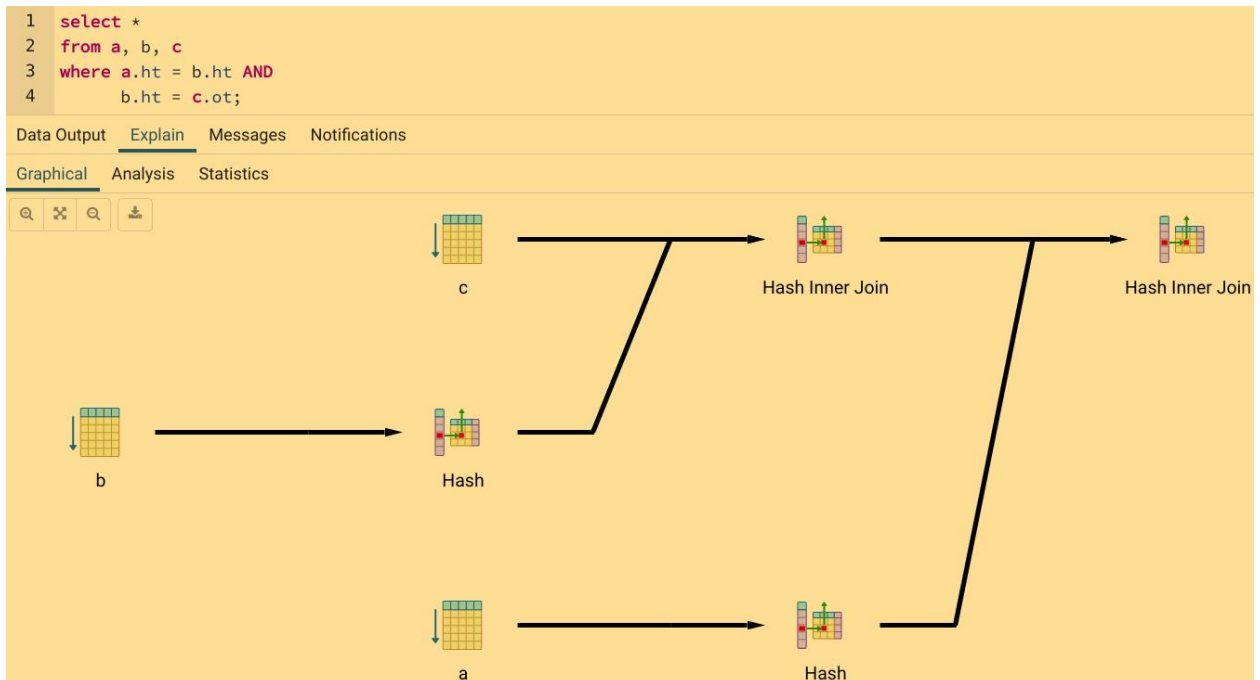


Things are different now -- this time it will first use merge join, together with index scans, to join the first two relations. Then use hash join to join the resulting table with the third table.

18.

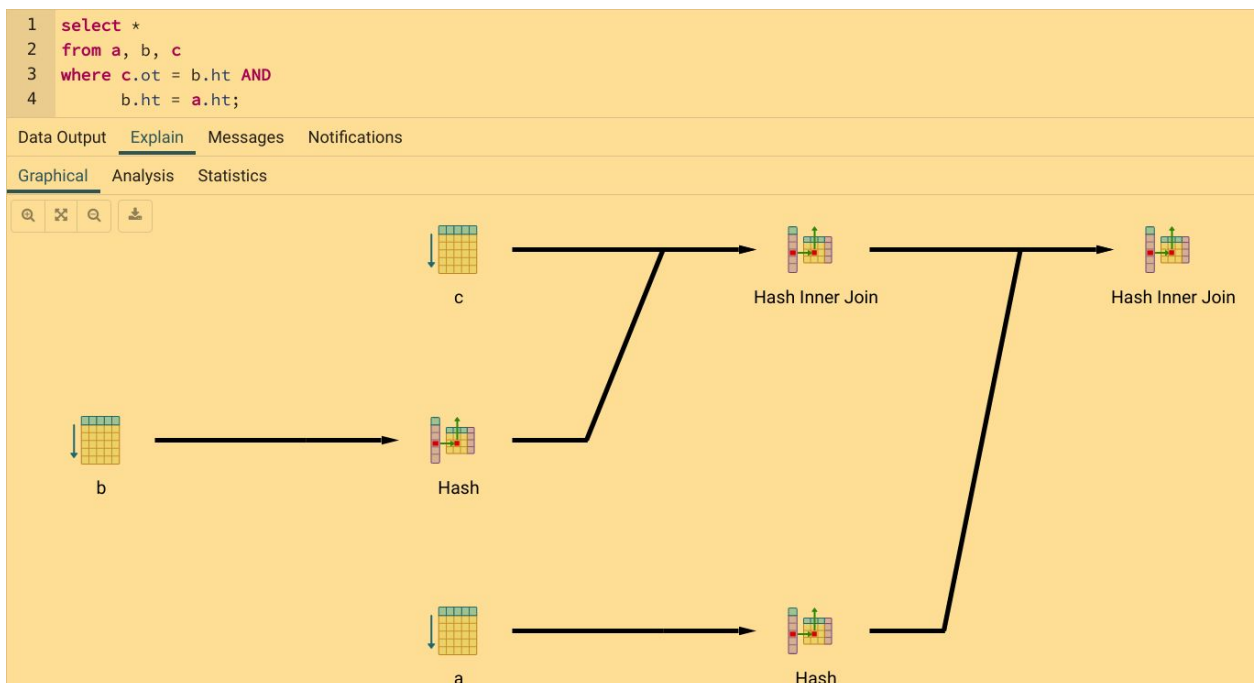
When joining on the "ten" or "hund" column, query plan 17 is chosen;
When joining on the "ot", "tt" or "ht" column, query plan 16 is chosen.

19.



Again, our most popular method -- hash join -- is chosen.

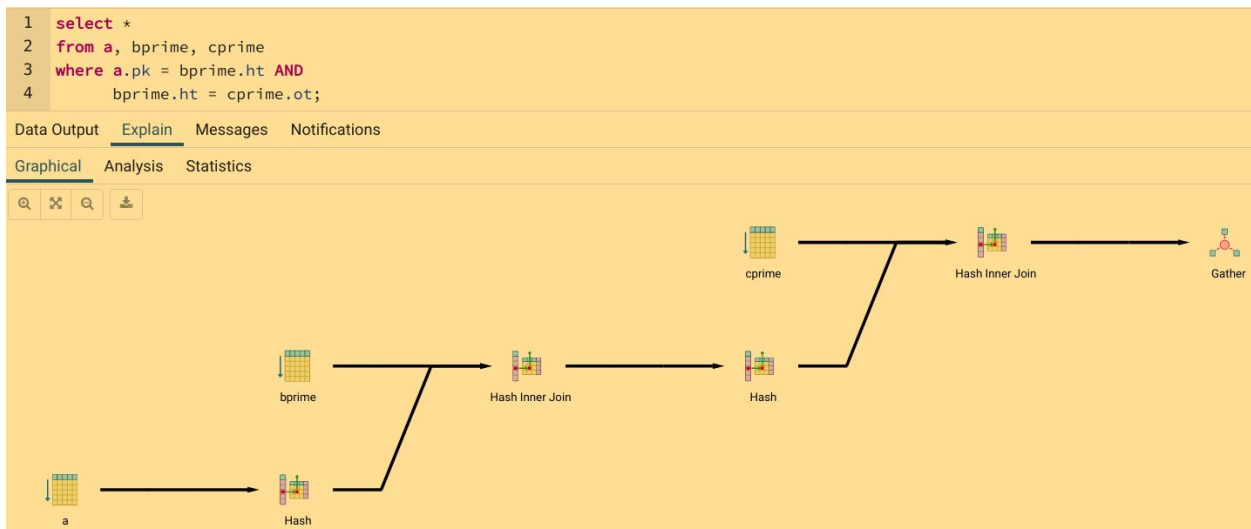
20.



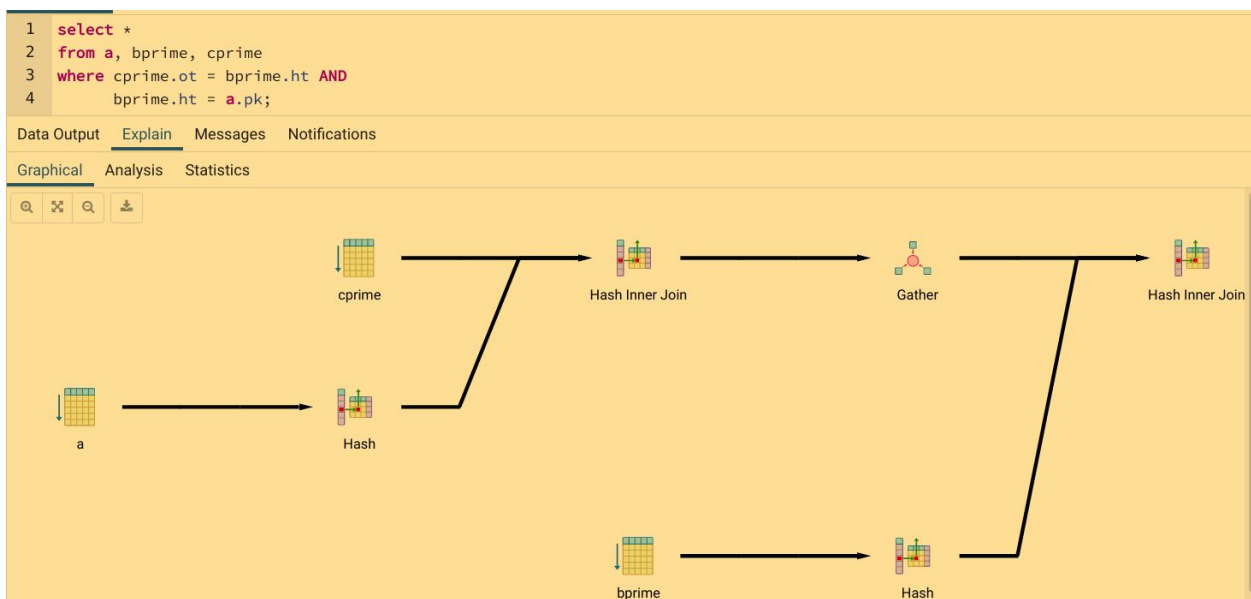
Exactly the same result as 19.

In 19 and 20, the result showed that the order of putting the relations in the query does not change the query plan chosen by the optimizer.

21.



22.



Hash join is used but the exact order has changed from 21.

In 21 and 22, the result showed that the actual algorithm chosen is not changed but the exact order of joining has changed. This is weird!

Part 3:

Review question 1:

Views. The SQL view construct has semantic roots in datalog. For example, suppose we have a relation `ElectronicDevice(product, model, price, batteryVolume)`. If we want to have a “laptop” view of it, we can do

Create view laptop as
Select model, price, batteryVolume
From ElectronicDevice
Where product = 'laptop';

Which can also be concisely expressed by the following Datalog semantics:

laptop(model, price, batteryVolume) <--
ElectronicDevice(product, model, price, batteryVolume) AND product = 'laptop'

Review question 2:
Enterprise Information Integration
Business logic/workflow

Review question 3:
 $\sigma_d(\text{Frequents} \bowtie \text{Likes} \bowtie \text{Sells})$

Review question 4:
The arity of Sells = 3
The arity of Happy = 1
The arity of Likes = 2

Review question 5:
Nondistinguished variables are "bar", "beer", and "p"

Review question 6:
There are 6 subgoals in the CheapBeer rule;
There is 1 distinguished variable.

Review question 7:
Assume each row in table Restaurant has a close time later than the open time.
Otherwise, don't input it into this relation in the first place.

(a) $\text{OpenDays}(r,d) \leftarrow \text{Restaurant}(r, d, _, _)$

(b) Define a new relation Days(d) that contains 7 rows in it -- "Sunday", "Monday", "Tuesday", "Wednesday", "Thursday", "Friday", "Saturday".
 $\text{ClosedDays}(r,d) \leftarrow \text{Days}(d) \text{ AND NOT } \text{OpenDays}(r, d)$

(c) No, because (UNOs 360, Sunday) is in ClosedDays.

(d) You might, because it's possible that UNOs 360 actually opens on Sundays but this information is not captured in the Restaurant relation. Hence the negated subgoal is not correctly computed.

Review question 8:

(a) Not safe because the distinguished variable "price" in the head does not show up anywhere in the body.

(b)

(c)